

▼ Animal Competition (15%)

Indented block

For the non-competition mode, we will use the Animal (<https://cloudstor.aarnet.edu.au/plus/s/cZYtNAeVhWD6uBX>) dataset. This dataset contains images of 151 different animals.

The dataset contains a total of 6270 images corresponding to the name of animal types.

All images are RGB images of 224 pixels wide by 224 pixels high in .jpg format. The images are separated in 151 folders according to their respective class.

The task is to categorize each animal into one of 151 categories.

We provide baseline code that includes the following features:

- Loading and Analysing the dataset using torchvision.
- Defining a simple convolutional neural network.
- How to use existing loss function for the model learning.
- Train the network on the training data.
- Test the trained network on the testing data.

The following changes could be considered:

1. "Transfer" Learning (ie use a model pre-trained another dataset)
2. Change of advanced training parameters: Learning Rate, Optimizer, Batch-size, Number of Max Epochs, and Drop-out.
3. Use of a new loss function.
4. Data augmentation
5. Architectural Changes: Batch Normalization, Residual layers, etc.
6. Others

Your code should be modified from the provided baseline. A pdf report of a maximum of two pages is required to explain the changes you made from the baseline, why you chose those changes, and the improvements they achieved.

Marking Rules:

We will mark the competition based on the final test accuracy on testing images and your report.

Final mark (out of 50) = acc_mark + efficiency mark + report mark

Acc_mark 10:

We will rank all the submission results based on their test accuracy. Zero improvement over the baseline yields 0 marks. Maximum improvement over the baseline will yield 10 marks. There will be a sliding scale applied in between.

Efficiency mark 10:

Efficiency considers not only the accuracy, but the computational cost of running the model (flops: <https://en.wikipedia.org/wiki/FLOPS>). Efficiency for our purposes is defined to be the ratio of accuracy (in %) to Gflops. Please report the computational cost for your final model and include the efficiency calculation in your report. Maximum improvement over the baseline will yield 10 marks. Zero improvement over the baseline yields zero marks, with a sliding scale in between.

Report mark 30:

Your report should comprise:

1. An introduction showing your understanding of the task and of the baseline model: [10 marks]
2. A description of how you have modified aspects of the system to improve performance. [10 marks]

A recommended way to present a summary of this is via an "ablation study" table, eg:

Method1	Method2	Method3	Accuracy
N	N	N	60%
Y	N	N	65%
Y	Y	N	77%
Y	Y	Y	82%

3. Explanation of the methods for reducing the computational cost and/or improve the trade-off between accuracy and cost: [5 marks]
4. Limitations/Conclusions: [5 marks]

```

1 #####
2 ### Subject: Computer Vision
3 ### Year: 2023
4 ### Student Name: Jian Yi Tai, Wei Long Wan
5 ### Student ID: a1836176, a1801549
6 ### Comptetion Name: Animal Classification Competition
7 ### Final Results:
8 ### ACC:          GFLOPs: 0.69
9 #####

```

```

1 # Importing libraries.
2
3 import os
4 import random
5 import numpy as np
6 import torch
7 import torch.nn as nn
8 import torch.nn.functional as F
9 from tqdm.notebook import tqdm
10
11 # To avoid non-essential warnings
12 import warnings
13 warnings.filterwarnings('ignore')
14
15 from torchvision import datasets, transforms, models
16 from torchvision.datasets import ImageFolder
17 from torchvision.transforms import ToTensor
18 from torchvision.utils import make_grid
19 from torch.utils.data import random_split
20 from torch.utils.data.dataloader import DataLoader
21 import matplotlib.pyplot as plt
22 %matplotlib inline

```

```

1 # Mounting G-Drive to get your dataset.
2 # To access Google Colab GPU; Go To: Edit >>> Netebook Settings >>> Hardware Accelarator: Select GPU.
3 # Reference: https://towardsdatascience.com/google-colab-import-and-export-datasets-eccf801e2971
4 from google.colab import drive
5 drive.mount('/content/drive')
6
7 # Dataset path. You should change the dataset path to the location that you place the data.
8 data_dir = '/content/drive/MyDrive/Colab Notebooks/Assignment4/animal/dataset/dataset'
9 classes = os.listdir(data_dir)

```

Mounted at /content/drive

```

1 # Performing Image Transformations.
2 ##Hints: Data Augmentation can be applied here. Have a look on RandomFlip, RandomRotation...
3 train_transform = transforms.Compose([
4     transforms.RandomResizedCrop(size=112, scale=(0.8, 1.0)), # changed from resize to random resize
5     transforms.RandomRotation(degrees=15), # added random rotation
6     transforms.RandomHorizontalFlip(),
7     transforms.CenterCrop(size=112),
8     transforms.ToTensor(),
9     transforms.Normalize((0.488), (0.2172))
10 ])

```

```

1 # Checking the dataset training size.
2 dataset = ImageFolder(data_dir, transform=train_transform)
3 print('Size of training dataset :', len(dataset))

```

Size of training dataset : 6270

```

1 # Viewing one of images shape.
2 img, label = dataset[100]
3 print(img.shape)

```

torch.Size([3, 112, 112])


```

1 # Preview one of the images..
2 def show_image(img, label):
3     print('Label: ', dataset.classes[label], "(" + str(label) + ")")
4     plt.imshow(img.permute(1,2,0))

```

show_image(*dataset[200])

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integ
Label: ailurus-fulgens (5)



```
1 # Setting seed so that value won't change everytime.
2 # Splitting the dataset to training, validation, and testing category.
3 torch.manual_seed(10)
4 val_size = len(dataset)//20
5 test_size = len(dataset)//10
6 train_size = len(dataset) - val_size - test_size

1 # Random Splitting.
2 train_ds, val_ds, test_ds = random_split(dataset, [train_size, val_size, test_size])
3 len(train_ds), len(val_ds), len(test_ds)

(5330, 313, 627)
```

```
1 batch_size = 16
2 train_loader = DataLoader(train_ds, batch_size, shuffle=True, num_workers=2, pin_memory=True)
3 val_loader = DataLoader(val_ds, batch_size, num_workers=2, pin_memory=True)
4 test_loader = DataLoader(test_ds, batch_size, num_workers=2, pin_memory=True)
```

```
1 # Multiple images preview.
2 for images, labels in train_loader:
3     fig, ax = plt.subplots(figsize=(18,10))
4     ax.set_xticks([])
5     ax.set_yticks([])
6     ax.imshow(make_grid(images, nrow=16).permute(1, 2, 0))
7     break
```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integ



```
1 # Baseline model class for training and validation purpose. Evaluation metric function - Accuracy.
2 def accuracy(output, target, topk=(1,)):
3     """
4     Computes the accuracy over the k top predictions for the specified values of k
5     In top-3 accuracy you give yourself credit for having the right answer
6     if the right answer appears in your top five guesses.
7     """
8     with torch.no_grad():
9         maxk = 3
10        batch_size = target.size(0)
11
12        # st()
13        _, pred = output.topk(maxk, 1, True, True)
14        pred = pred.t()
15        # st()
16        # correct = pred.eq(target.view(1, -1).expand_as(pred))
17        # correct = (pred == target.view(1, -1).expand_as(pred))
18        correct = (pred == target.unsqueeze(dim=0)).expand_as(pred)
19
20
21
22        correct_3 = correct[:3].reshape(-1).float().sum(0, keepdim=True)
23
```

```

24         return correct_3.mul_(1.0 / batch_size)
25 #def accuracy(outputs, labels):
26 #     _, preds = torch.max(outputs, dim=1)
27 #     return torch.tensor(torch.sum(preds == labels).item() / len(preds))
28
29 class ImageClassificationBase(nn.Module):
30     def training_step(self, batch):
31         images, labels = batch
32         out = self(images)             # Generate predictions
33         loss = F.cross_entropy(out, labels)
34         return loss
35
36     def validation_step(self, batch):
37         images, labels = batch
38         out = self(images)             # Generate predictions
39         loss = F.cross_entropy(out, labels)
40         acc = accuracy(out, labels, (5))    # Calculate accuracy
41         return {'val_loss': loss.detach(), 'val_acc': acc}
42
43     def validation_epoch_end(self, outputs):
44         batch_losses = [x['val_loss'] for x in outputs]
45         epoch_loss = torch.stack(batch_losses).mean()    # Combine losses
46         batch_accs = [x['val_acc'] for x in outputs]
47         epoch_acc = torch.stack(batch_accs).mean()      # Combine accuracies
48         return {'val_loss': epoch_loss.item(), 'val_acc': epoch_acc.item()}
49
50     def epoch_end(self, epoch, result):
51         print("Epoch [{}], train_loss: {:.4f}, val_loss: {:.4f}, val_acc: {:.4f}".format(
52             epoch, result['train_loss'], result['val_loss'], result['val_acc']))

```

```

1 # To check whether Google Colab GPU has been assigned/not.

```

```

2
3 def get_default_device():
4     """Pick GPU if available, else CPU"""
5     if torch.cuda.is_available():
6         return torch.device('cuda')
7     else:
8         return None
9
10 def to_device(data, device):
11     """Move tensor(s) to chosen device"""
12     if isinstance(data, (list,tuple)):
13         return [to_device(x, device) for x in data]
14     return data.to(device, non_blocking=True)
15
16 class DeviceDataLoader():
17     """Wrap a dataloader to move data to a device"""
18     def __init__(self, dl, device):
19         self.dl = dl
20         self.device = device
21
22     def __iter__(self):
23         """Yield a batch of data after moving it to device"""
24         for b in self.dl:
25             yield to_device(b, self.device)
26
27     def __len__(self):
28         """Number of batches"""
29         return len(self.dl)

```

```

1 device = get_default_device()
2 device
3 train_loader = DeviceDataLoader(train_loader, device)
4 val_loader = DeviceDataLoader(val_loader, device)
5 test_loader = DeviceDataLoader(test_loader, device)

```

```

1 input_size = 3*112*112
2 output_size = 151

```

```

1 # Convolutional Network - Baseline
2 class ConvolutionalNetwork(ImageClassificationBase):
3     def __init__(self, classes):
4         super().__init__()
5         self.num_classes=classes
6         self.conv1=nn.Conv2d(3,64,5,1)
7         self.conv1_bn=nn.BatchNorm2d(64)    # normalize the layers
8         self.conv2=nn.Conv2d(64,128,3,1)
9         self.conv2_bn=nn.BatchNorm2d(128)

```

```

10     self.conv3=nn.Conv2d(128,128,3,1)
11     self.conv3_bn=nn.BatchNorm2d(128)
12     self.conv4=nn.Conv2d(128,128,3,1)
13     self.conv4_bn=nn.BatchNorm2d(128)
14     self.fc1=nn.Linear(128*5*5,self.num_classes)
15
16     def forward(self,X):
17
18         X=self.conv1(X)
19         X=F.relu(self.conv1_bn(X))
20         X=F.max_pool2d(X,2,2)
21         X=self.conv2(X)
22         X=F.relu(self.conv2_bn(X))
23         X=F.max_pool2d(X,2,2)
24         X=self.conv3(X)
25         X=F.relu(self.conv3_bn(X))
26         X=F.max_pool2d(X,2,2)
27         X=self.conv4(X)
28         X=F.relu(self.conv4_bn(X))
29         X=F.max_pool2d(X,2,2)
30
31         X=X.view(-1,128*5*5)
32         X=self.fc1(X)
33
34         return F.log_softmax(X, dim=1)

1 # Model print
2 num_classes = 151
3 model = ConvolutionalNetwork(num_classes)
4 model.cuda()

ConvolutionalNetwork(
  (conv1): Conv2d(3, 64, kernel_size=(5, 5), stride=(1, 1))
  (conv1_bn): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1))
  (conv2_bn): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1))
  (conv3_bn): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv4): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1))
  (conv4_bn): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (fc1): Linear(in_features=3200, out_features=151, bias=True)
)

1 # We can check the input and the output shape
2 for images, labels in train_loader:
3     images = images.cuda()
4     out = model(images)
5     print('images.shape:', images.shape)
6     print('out.shape:', out.shape)
7     print('out[0]:', out[0])
8     break

images.shape: torch.Size([16, 3, 112, 112])
out.shape: torch.Size([16, 151])
out[0]: tensor([-5.3641, -5.6836, -5.5946, -5.7739, -6.1157, -4.9951, -4.1743, -4.7631,
-5.7735, -4.9621, -4.9313, -5.9522, -5.2435, -5.4704, -5.6282, -5.1503,
-4.5002, -5.8129, -5.6460, -5.3197, -6.3240, -5.5163, -5.8242, -5.0268,
-5.9893, -5.1987, -4.5718, -5.5512, -5.4439, -5.4620, -4.8981, -4.6023,
-6.0071, -5.0184, -4.5394, -5.7638, -6.0071, -4.2020, -4.2432, -4.4801,
-5.4682, -6.1645, -4.7778, -5.2560, -4.2193, -5.2059, -5.1022, -5.5018,
-5.8577, -4.8406, -5.1738, -5.7383, -4.5258, -6.2852, -5.7984, -5.8892,
-5.6387, -5.3767, -5.2718, -4.5317, -5.2374, -4.7273, -4.7446, -4.8545,
-5.5233, -4.7681, -5.1486, -4.6807, -4.2106, -5.9732, -5.2196, -4.7531,
-5.6543, -4.6611, -6.5081, -5.2076, -5.2105, -4.9096, -5.4934, -4.6193,
-6.2550, -5.6043, -4.6507, -4.4868, -5.2780, -4.5580, -4.9255, -5.6942,
-6.1539, -5.4724, -4.9867, -5.2039, -4.6595, -4.0290, -5.2641, -5.2490,
-4.6138, -5.1483, -5.8443, -5.1778, -5.4398, -4.7622, -5.0176, -4.4051,
-4.5823, -4.8575, -4.8328, -5.6341, -4.8259, -5.2256, -5.1356, -5.5450,
-5.5738, -5.8489, -5.3093, -5.1925, -5.2229, -4.5316, -4.6551, -6.1339,
-4.9697, -6.2367, -5.2043, -4.8287, -5.1412, -4.4861, -4.4355, -4.8280,
-3.4915, -4.2773, -5.0159, -5.1184, -5.2187, -5.2487, -4.9208, -4.4919,
-5.3626, -5.9790, -4.1744, -4.9083, -5.6280, -5.5590, -5.9117, -4.7588,
-4.6044, -5.9651, -4.2713, -5.4493, -5.5173, -5.4419, -4.1202],
device='cuda:0', grad_fn=<SelectBackward0>)

1 train_dl = DeviceDataLoader(train_loader, device)
2 val_dl = DeviceDataLoader(val_loader, device)
3 to_device(model, device)

ConvolutionalNetwork(
  (conv1): Conv2d(3, 64, kernel_size=(5, 5), stride=(1, 1))
  (conv1_bn): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

```

(conv2): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1))
(conv2_bn): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(conv3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1))
(conv3_bn): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(conv4): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1))
(conv4_bn): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(fc1): Linear(in_features=3200, out_features=151, bias=True)
)

1 #https://stackoverflow.com/questions/71998978/early-stopping-in-pytorch
2 ### Early stopper ###
3 class EarlyStopper:
4     def __init__(self, patience=1, min_delta=0):
5         self.patience = patience
6         self.min_delta = min_delta
7         self.counter = 0
8         self.min_validation_loss = np.inf
9
10    def early_stop(self, validation_loss):
11        if validation_loss < self.min_validation_loss:
12            self.min_validation_loss = validation_loss
13            self.counter = 0
14        elif validation_loss > (self.min_validation_loss + self.min_delta):
15            self.counter += 1
16            if self.counter >= self.patience:
17                return True
18        return False

1 from torch.optim import lr_scheduler
2 # Functions for evaluation and training.
3
4 @torch.no_grad()
5 def evaluate(model, val_loader):
6     model.eval()
7     outputs = [model.validation_step(batch) for batch in val_loader]
8     return model.validation_epoch_end(outputs)
9
10 def fit(patience, min_delta, epochs, lr, model, train_loader, val_loader, opt_func=torch.optim.SGD):
11     history = []
12     optimizer = opt_func(model.parameters(), lr)
13     scheduler = lr_scheduler.ExponentialLR(optimizer, gamma=0.9) # scheduler
14     early_stopper = EarlyStopper(patience, min_delta)
15     for epoch in range(epochs):
16         # Training Phase
17         model.train()
18         train_losses = []
19         for batch in tqdm(train_loader):
20             loss = model.training_step(batch)
21             train_losses.append(loss)
22             loss.backward()
23             optimizer.step()
24             optimizer.zero_grad()
25         scheduler.step()
26
27         # Validation phase
28         result = evaluate(model, val_loader)
29         result['train_loss'] = torch.stack(train_losses).mean().item()
30         model.epoch_end(epoch, result)
31         history.append(result)
32         if early_stopper.early_stop(result['val_loss']):
33             break
34     return history

1 model = to_device(model, device)

1 history=[evaluate(model, val_loader)]
2 history

[{'val_loss': 5.015705108642578, 'val_acc': 0.01875000074505806}]

1 # Hints: The following parameters can be changed to improve the accuracy
2 print(test_size)
3 num_epochs = 100
4 opt_func = torch.optim.Adam
5 lr = 0.001

```

```
1 # del model
2 # del history
```

```
1 # fit (patience, min_delta)
2 history+= fit(5, 0, num_epochs, lr, model, train_dl, val_dl, opt_func)
```

```
100% 334/334 [13:22<00:00, 2.42s/it]
Epoch [0], train_loss: 4.8812, val_loss: 4.1871, val_acc: 0.1823
100% 334/334 [00:21<00:00, 17.65it/s]
Epoch [1], train_loss: 3.9772, val_loss: 3.8540, val_acc: 0.2972
100% 334/334 [00:20<00:00, 20.77it/s]
Epoch [2], train_loss: 3.5273, val_loss: 3.6729, val_acc: 0.3042
100% 334/334 [00:21<00:00, 18.19it/s]
Epoch [3], train_loss: 3.1364, val_loss: 3.4364, val_acc: 0.3753
100% 334/334 [00:19<00:00, 20.57it/s]
Epoch [4], train_loss: 2.8542, val_loss: 3.3192, val_acc: 0.4021
100% 334/334 [00:21<00:00, 20.76it/s]
Epoch [5], train_loss: 2.6167, val_loss: 3.2621, val_acc: 0.4365
100% 334/334 [00:19<00:00, 21.28it/s]
Epoch [6], train_loss: 2.3037, val_loss: 3.2326, val_acc: 0.4615
100% 334/334 [00:21<00:00, 16.58it/s]
Epoch [7], train_loss: 2.0847, val_loss: 3.0424, val_acc: 0.4865
100% 334/334 [00:19<00:00, 21.22it/s]
Epoch [8], train_loss: 1.8933, val_loss: 3.1625, val_acc: 0.4708
100% 334/334 [00:21<00:00, 16.95it/s]
Epoch [9], train_loss: 1.7068, val_loss: 3.0602, val_acc: 0.5264
100% 334/334 [00:19<00:00, 18.21it/s]
Epoch [10], train_loss: 1.5221, val_loss: 2.9845, val_acc: 0.5389
100% 334/334 [00:20<00:00, 15.12it/s]
Epoch [11], train_loss: 1.3375, val_loss: 3.0149, val_acc: 0.5208
100% 334/334 [00:18<00:00, 19.25it/s]
Epoch [12], train_loss: 1.2264, val_loss: 2.9717, val_acc: 0.5139
100% 334/334 [00:18<00:00, 17.17it/s]
Epoch [13], train_loss: 1.1051, val_loss: 3.0646, val_acc: 0.5358
100% 334/334 [00:19<00:00, 20.51it/s]
Epoch [14], train_loss: 1.0008, val_loss: 2.9578, val_acc: 0.5771
100% 334/334 [00:18<00:00, 21.37it/s]
Epoch [15], train_loss: 0.9331, val_loss: 3.1263, val_acc: 0.5538
100% 334/334 [00:20<00:00, 19.14it/s]
Epoch [16], train_loss: 0.8548, val_loss: 3.0088, val_acc: 0.5521
100% 334/334 [00:18<00:00, 20.71it/s]
Epoch [17], train_loss: 0.7951, val_loss: 3.0340, val_acc: 0.5521
100% 334/334 [00:19<00:00, 12.84it/s]
Epoch [18], train_loss: 0.7399, val_loss: 3.0565, val_acc: 0.5576
100% 334/334 [00:18<00:00, 20.97it/s]
Epoch [19], train_loss: 0.6903, val_loss: 2.9359, val_acc: 0.5708
100% 334/334 [00:18<00:00, 21.67it/s]
Epoch [20], train_loss: 0.6751, val_loss: 2.9025, val_acc: 0.5851
100% 334/334 [00:19<00:00, 20.43it/s]
Epoch [21], train_loss: 0.6054, val_loss: 2.9035, val_acc: 0.5771
100% 334/334 [00:18<00:00, 20.48it/s]
Epoch [22], train_loss: 0.5845, val_loss: 3.0867, val_acc: 0.5253
100% 334/334 [00:19<00:00, 13.69it/s]
Epoch [23], train_loss: 0.5716, val_loss: 2.9819, val_acc: 0.5528
100% 334/334 [00:18<00:00, 17.85it/s]
Epoch [24], train_loss: 0.5404, val_loss: 3.0120, val_acc: 0.5753
100% 334/334 [00:18<00:00, 19.16it/s]
Epoch [25], train_loss: 0.5099, val_loss: 2.9412, val_acc: 0.6021
```

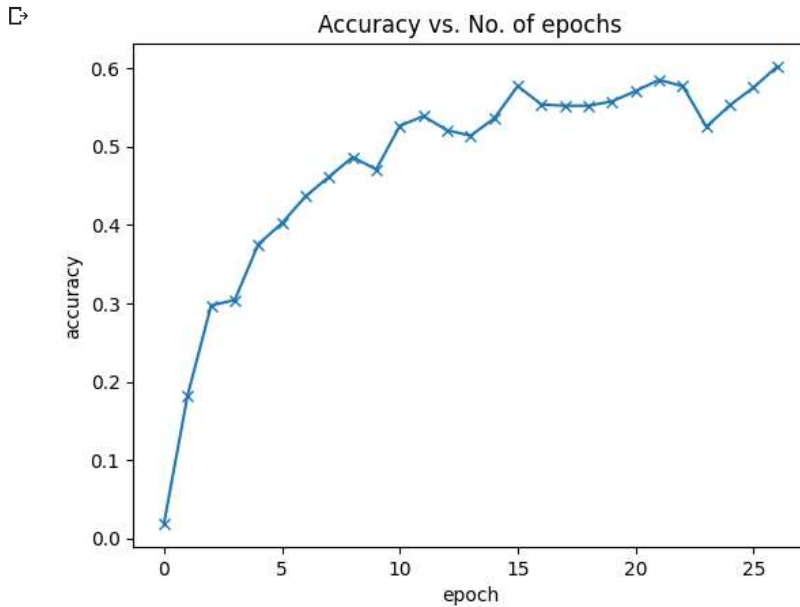
```
1 def plot_accuracies(history):
2     accuracies = [x['val_acc'] for x in history]
3     plt.plot(accuracies, '-x')
4     plt.xlabel('epoch')
5     plt.ylabel('accuracy')
6     plt.title('Accuracy vs. No. of epochs')
```

```

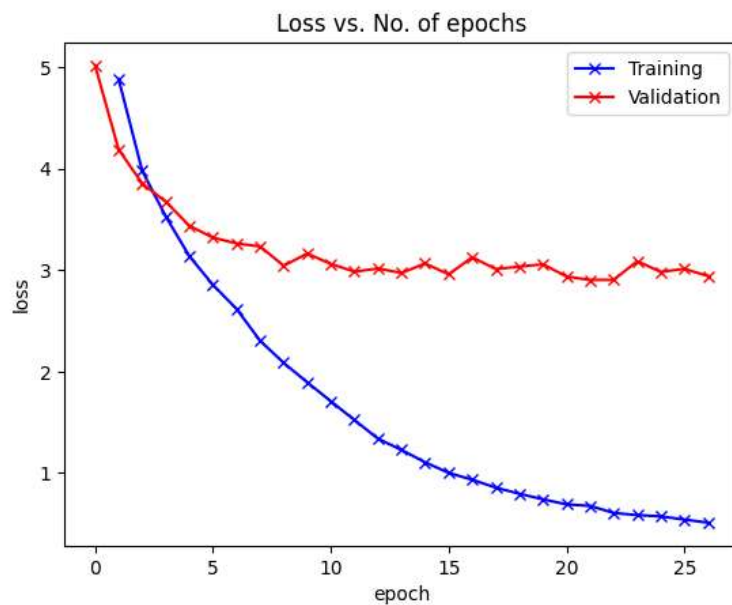
6 plt.title('Accuracy vs. No. of epochs')
7 plt.show()
8
9 def plot_losses(history):
10     train_losses = [x.get('train_loss') for x in history]
11     val_losses = [x['val_loss'] for x in history]
12     plt.plot(train_losses, '-bx')
13     plt.plot(val_losses, '-rx')
14     plt.xlabel('epoch')
15     plt.ylabel('loss')
16     plt.legend(['Training', 'Validation'])
17     plt.title('Loss vs. No. of epochs')
18     plt.show()

```

```
1 plot_accuracies(history)
```



```
1 plot_losses(history)
```



```
1 evaluate(model, test_loader)
```

```
{'val_loss': 2.845857858657837, 'val_acc': 0.6348958611488342}
```

▼ FLOPs

```

1 #The code from https://cloudstor.aarnet.edu.au/plus/s/PcSc67ZncTSQP0E can be used to count flops
2 #Download the code.
3 !wget -c https://cloudstor.aarnet.edu.au/plus/s/hXo1dK9SZqiEVn9/download

```



```
4 !mv download FLOPs_counter.py
5 #!rm -rf download
```

```
--2023-06-22 12:05:55-- https://cloudstor.aarnet.edu.au/plus/s/hXo1dk9SZqiEVn9/download
Resolving cloudstor.aarnet.edu.au (cloudstor.aarnet.edu.au)... 202.158.207.20
Connecting to cloudstor.aarnet.edu.au (cloudstor.aarnet.edu.au)|202.158.207.20|:443... connected.
HTTP request sent, awaiting response... 200 OK
Syntax error in Set-Cookie: 5230042dc1897=1j39fk7h7nd8v7kcs999cu42kh; path=/plus; domain=.aarnet.edu.au;; Secure; SameSite=Lax at
Syntax error in Set-Cookie: oc_sessionPassphrase=zyV0F8FdY0s1bmPF2px%2Fcws9x4EcKcdBwPn1oV4q90up1Ccn0o4NRNzuz%2BNOrRL9tI8TnBXZYtU
Length: 5201 (5.1K) [text/x-python]
Saving to: 'download'
```

```
download          100%[=====] 5.08K --.-KB/s in 0s
```

```
2023-06-22 12:05:56 (756 MB/s) - 'download' saved [5201/5201]
```



```
1 from FLOPs_counter import print_model_parm_flops
2 input = torch.randn(1, 3, 112, 112) # The input size should be the same as the size that you put into your model
3 #Get the network and its FLOPs
4 num_classes = 151
5 model = ConvolutionalNetwork(num_classes)
6 print_model_parm_flops(model, input, detail=False)

+ Number of FLOPs: 0.69G
```

✓ 0s completed at 9:36 PM

